

Internet Protocols EBU5403

Michael Chai (michael.chai@qmul.ac.uk)

Richard Clegg (r.clegg@qmul.ac.uk)

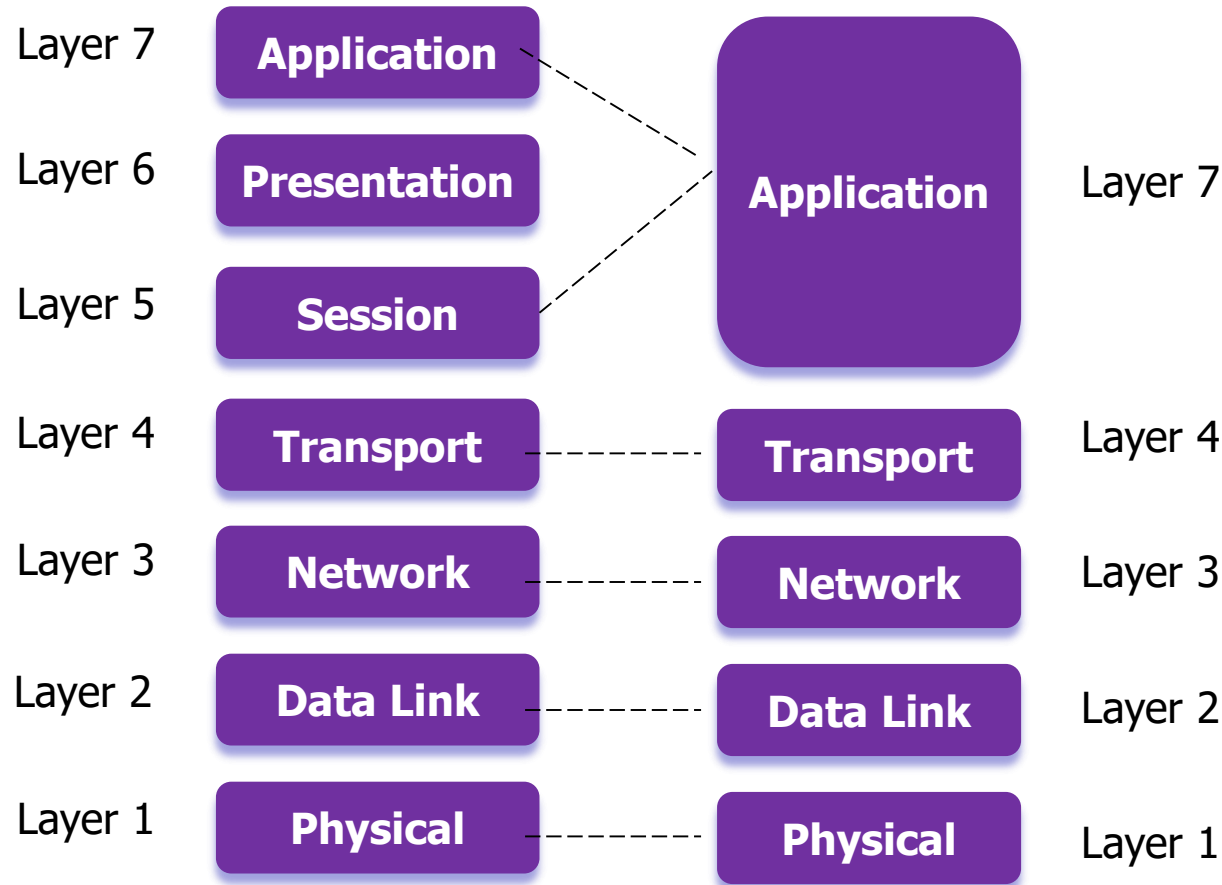
Cunhua Pan (c.pan@qmul.ac.uk)

	Week 1	Week 2	Week 3	Week 4
Group 1	Michael	Cunhua	Michael	Cunhua
Group 2	Richard			
Group 3	Michael	Cunhua	Michael	Cunhua

Structure of course

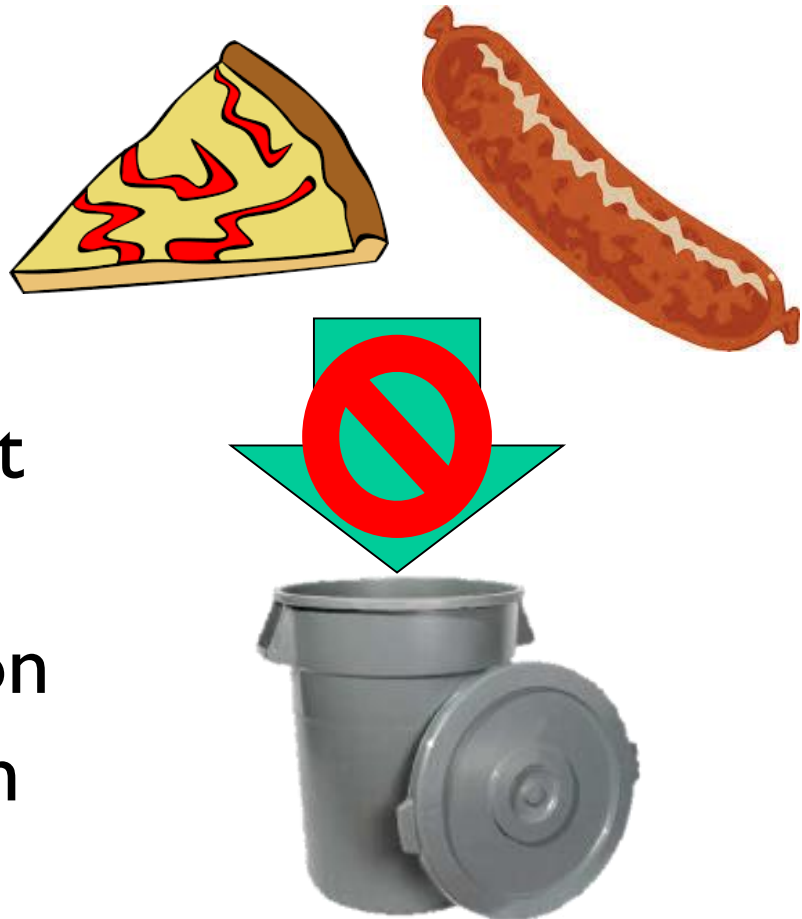
- Week 1
 - Introduction to IP Networks
 - The Transport layer (part I)
- Week 2
 - The Transport layer (part II)
 - The Network layer (part I)
 - Class test
- Week 3
 - The Network layer (part II)
 - The Data link layer (part I)
 - Router lab tutorial (assessed lab work after this week)
- Week 4
 - The Data link layer (part II)
 - Network management and security
 - Class test

ISO/OSI (left) vs TCP/IP (right)



How to remember the layers

- Please Do Not Throw Pizza + Sausage Away
- **P**lease – **P**hysical
- **D**o – **D**atalink
- **N**ot – **N**etwork
- **T**hrow – **T**ransport
- **S**ausage – **S**ession
- **P**izza – **P**resentation
- **A**way – **A**pplication

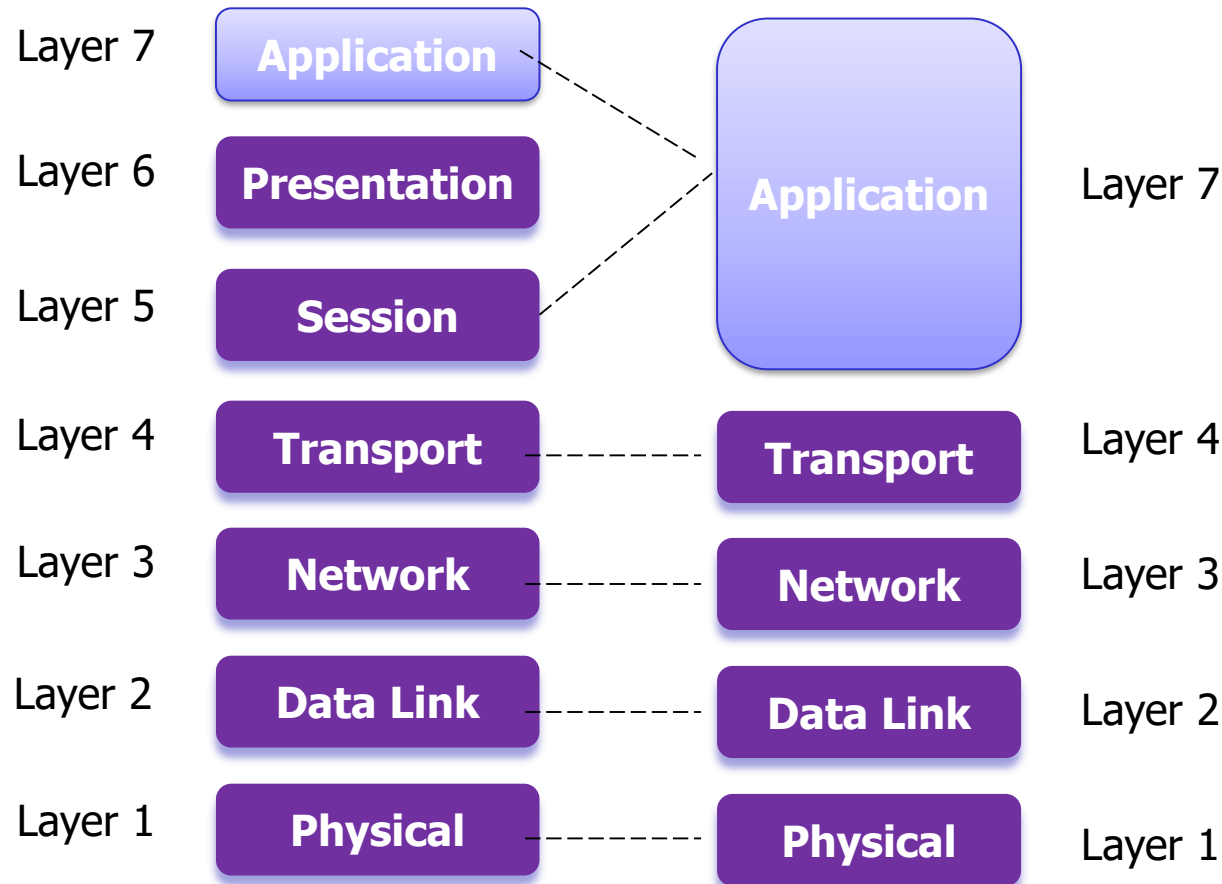


Week 1: Transport Layer (part 1)

our goals:

- understand principles behind transport layer services:
 - multiplexing, demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
- learn about Internet transport layer protocols:
 - UDP: connectionless transport
 - TCP: connection-oriented reliable transport
 - TCP congestion control

Application layer (very briefly)

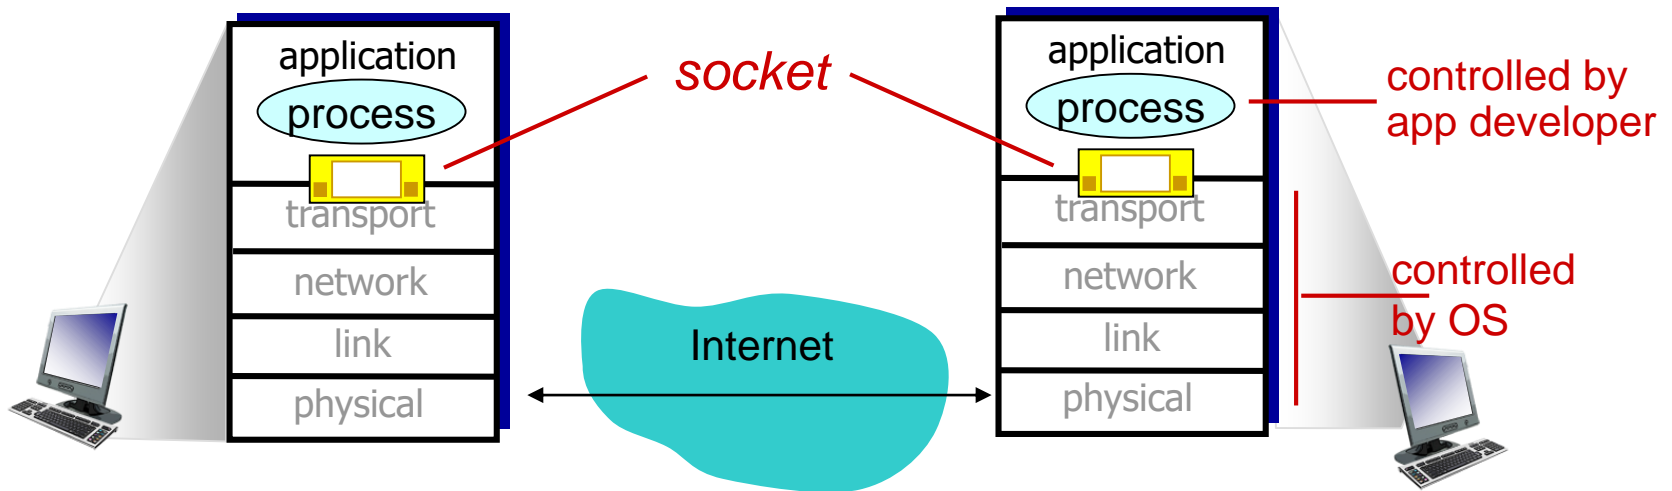


Application Layer (very briefly)

- This is covered in many other modules you will take.
- This is the layer you (mostly) work with as a programmer.
- At the application layer the “network” is abstracted away and you access a stream of data that arrives at a “socket”.
- It is up to you to define the format of data your application writes and receives.
- Different applications have different formats:
 - HTTP (Hypertext transfer protocol) world-wide web
 - FTP (File transfer protocol) moving data
 - SMTP (Send Mail transfer protocol) sending email
 - IMAP (Internet message access protocol) receiving email

Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



Addressing processes

- to receive messages, process must have *identifier*
- host device has unique 32-bit IP address
- Q: does IP address of host on which process runs suffice for identifying the process?
 - A: no, *many* processes can be running on same host
- *identifier* includes both **IP address** and **port numbers** associated with process on host.
- example port numbers:
 - HTTP server: 80
 - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
 - **IP address**: 128.119.245.12
 - **port number**: 80
- more shortly...

4 different addresses in TCP/IP

- Physical address Layer 2
 - Also known as the link address, is the address of a node as defined by its LAN or WAN
- Logical address Layer 3 (32-bit , IPv4 128-bit IPv6)
 - Logical (IP) addresses are for universal communications that are independent of underlying physical networks
- Port address Layer 4 (16-bit)
 - Port addresses differentiate different processes
- Application-specific address Layer 7
 - Some applications have user-friendly addresses that are designed for that specific application, such as email address, URL.

App-layer protocol defines

- **types of messages exchanged,**
 - e.g., request, response
- **message syntax:**
 - what fields in messages & how fields are delineated
- **message semantics**
 - meaning of information in fields
- **rules** for when and how processes send & respond to messages

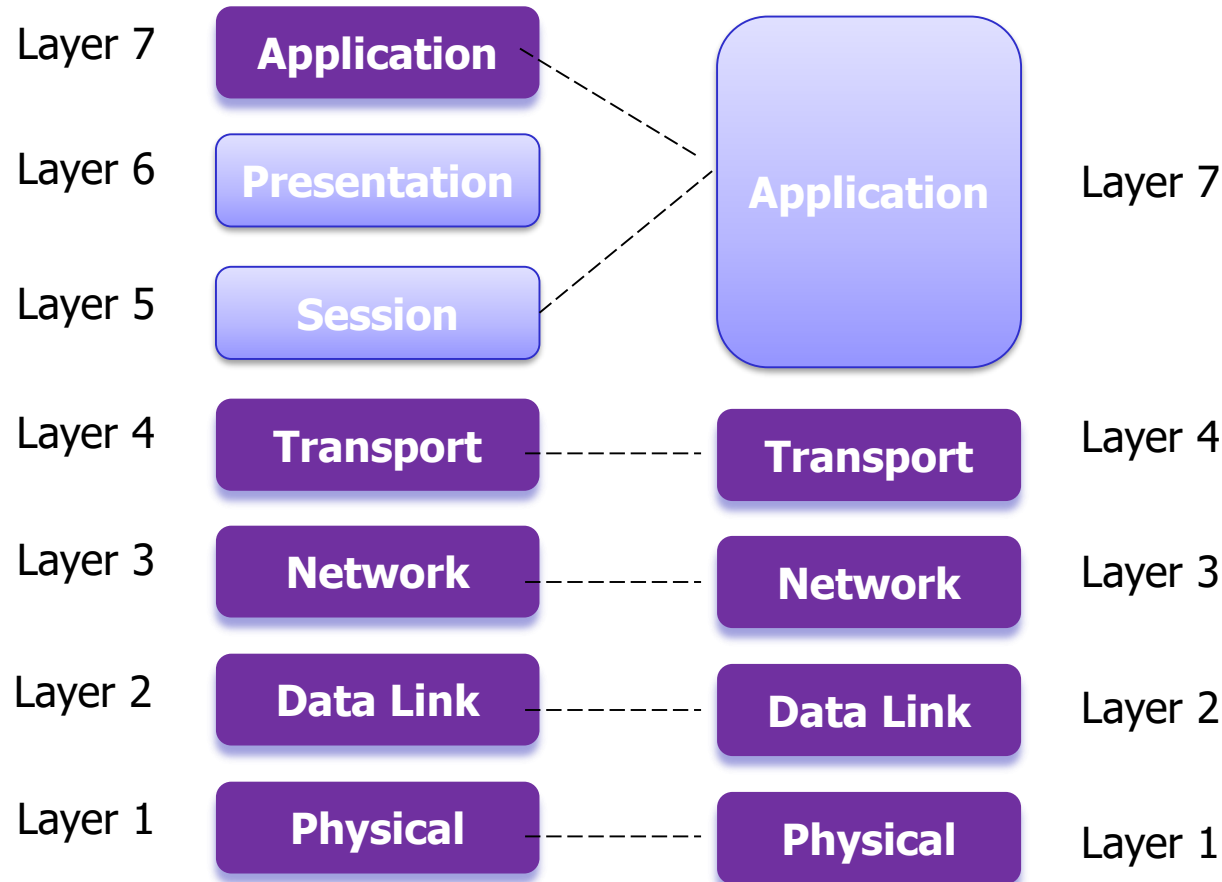
open protocols:

- defined in RFCs (request for comments)
- allows for interoperability
- e.g., HTTP (web), SMTP (email)

proprietary protocols:

- e.g., Skype

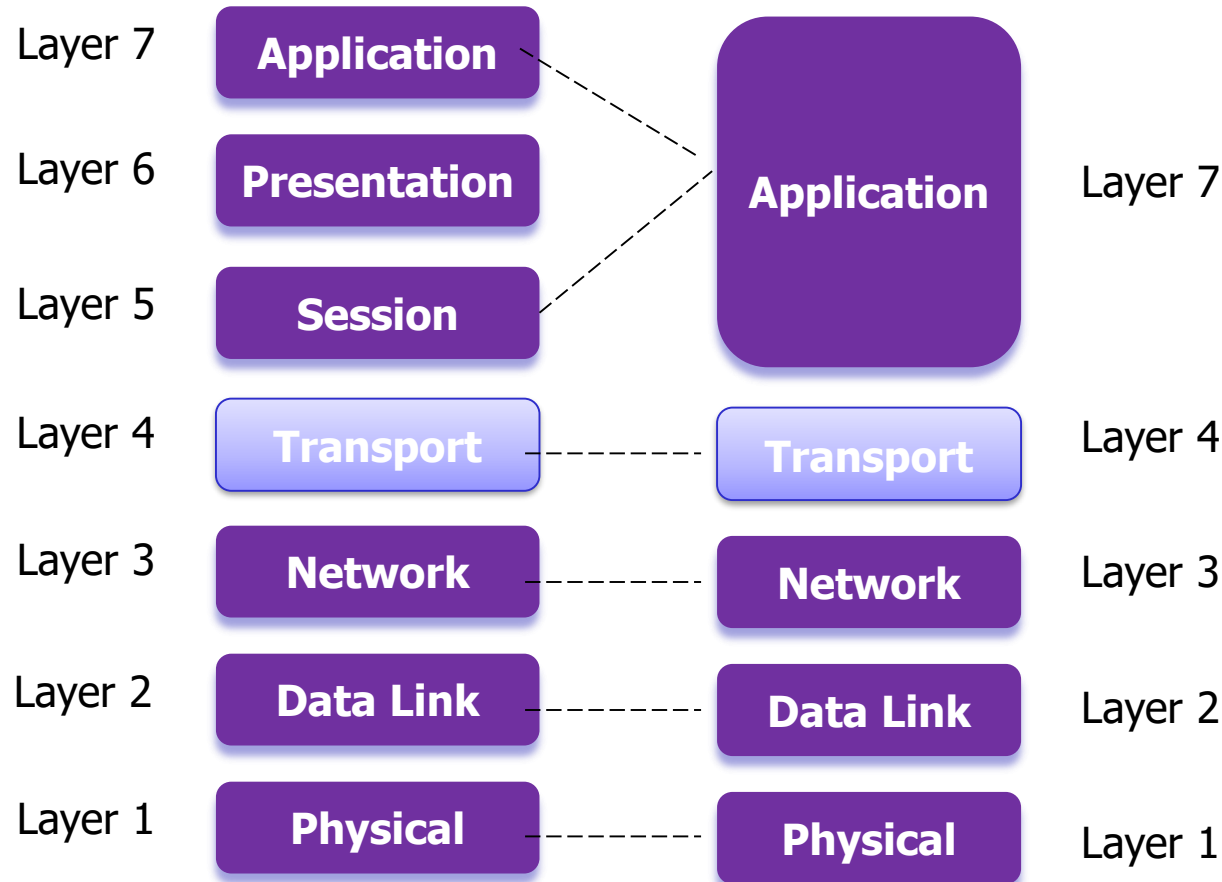
Session + Presentation layers



Session and Presentation Layers

- Theoretical only – not implemented in the Internet
- Session Layer (layer 5):
 - Takes care of a connection between two hosts for the lifetime of that connection
 - Authentication + authorisation (who is the user, what can they do)
 - In working Internet this is at the application layer.
- Presentation Layer (layer 6):
 - Translates data for an application.
 - For example takes care of detail of character set used to encode string of characters.
 - In working Internet this is at the application layer.

Transport Layer



Transport layer: outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

Part I (this week)

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Part II (next week)

Numbers relate to chapter numbers in
Course text: Kurose + Ross

Transport Layer: outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

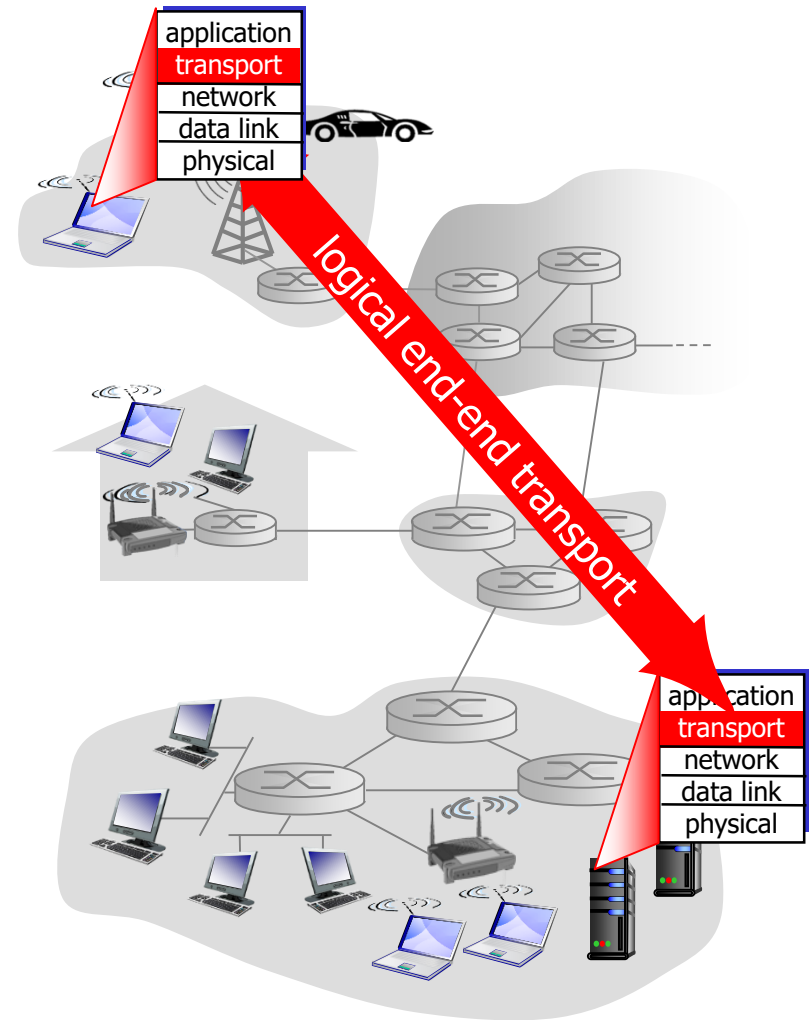
3.7 TCP congestion control

How to remember the layers

- Please Do Not Throw Pizza + Sausage Away
- **P**lease – **P**hysical
- **D**o – **D**atalink
- **N**ot – **N**etwork
- **T**hrow – **T**ransport
- **S**ausage – **S**ession
- **P**izza – **P**resentation
- **A**way – **A**pplication

Transport services and protocols

- provide *logical communication* between app processes running on different hosts
- transport protocols run in end systems
 - send side: breaks app messages into *segments*, passes to network layer
 - rcv side: reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
 - Internet: TCP and UDP



Transport vs. network layer

- *network layer*: logical communication between hosts
- *transport layer*: logical communication between processes
 - relies on, enhances, network layer services

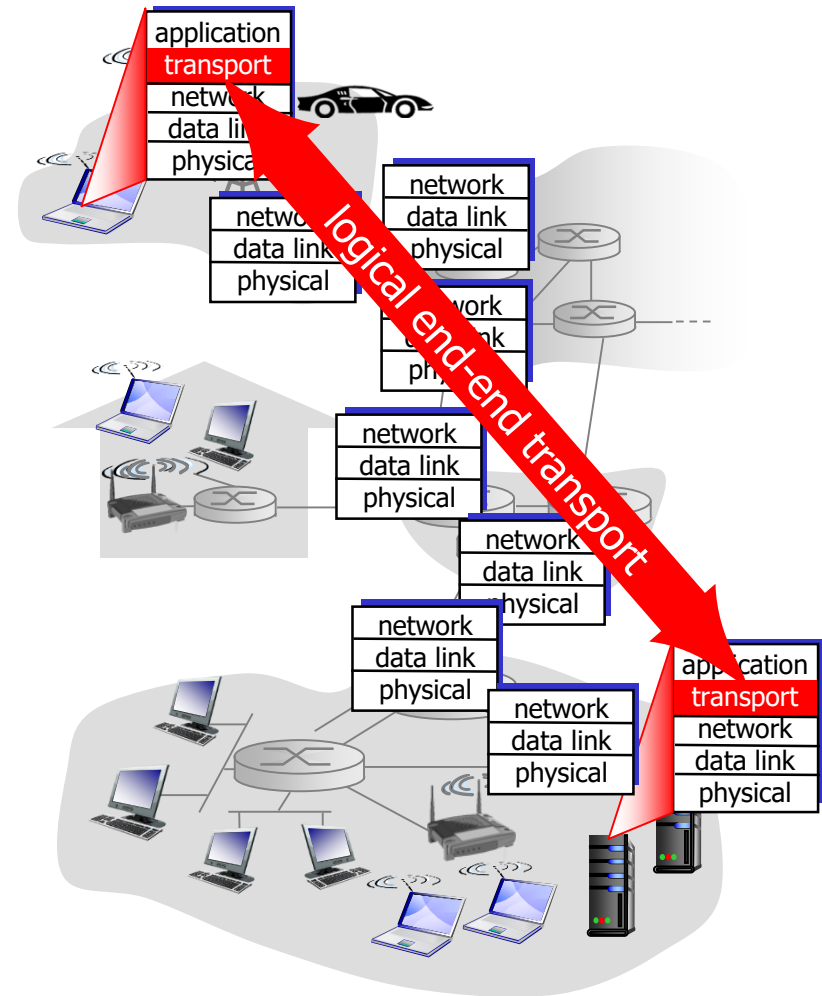
household analogy:

12 kids in Ann's house sending letters to 12 children in Bill's house:

- hosts = houses
- processes = children
- app messages = letters in envelopes
- transport protocol = Ann and Bill who give letter to correct child
- network-layer protocol = postal service

Internet transport-layer protocols

- reliable, in-order delivery (TCP)
 - congestion control
 - flow control
 - connection setup
- unreliable, unordered delivery: UDP
 - no-frills extension of “best-effort” IP
- services not available:
 - delay guarantees
 - bandwidth guarantees



Transport Layer: outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Multiplexing/demultiplexing

■ Multiplexing (Mux):

- Combining several streams of data into a single stream.
- Example – your phone is browsing the web, refreshing your email, connecting to wechat at the same time. All these connections are sent over the same link.

■ Demultiplexing (Demux)

- The opposite process. A stream of data is separated out into its individual components.
- The stream of packets the phone received is split up and sent to the appropriate program for web, email, wechat.

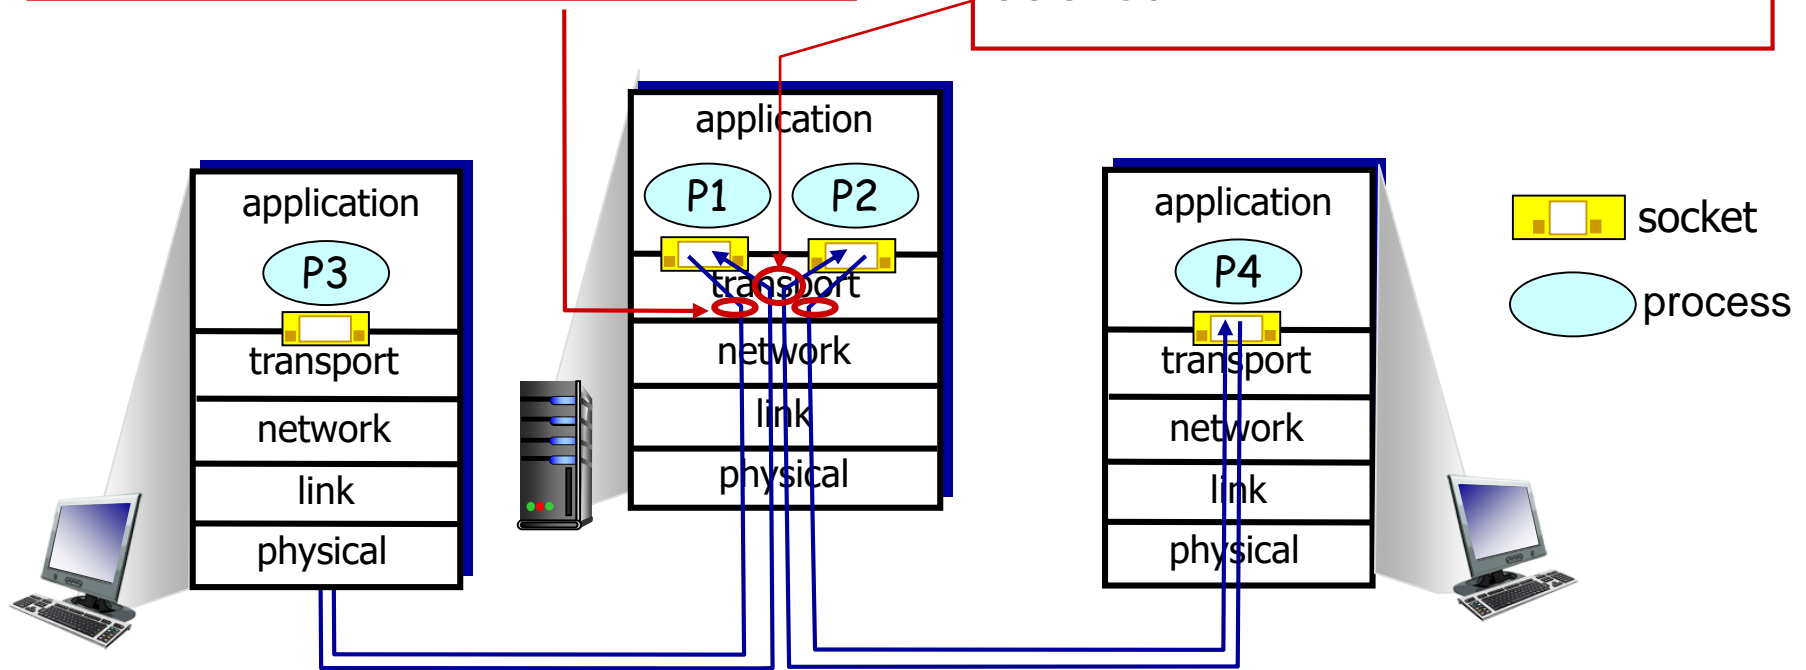
Multiplexing/demultiplexing

multiplexing at sender:

handle data from multiple sockets, add transport header (later used for demultiplexing)

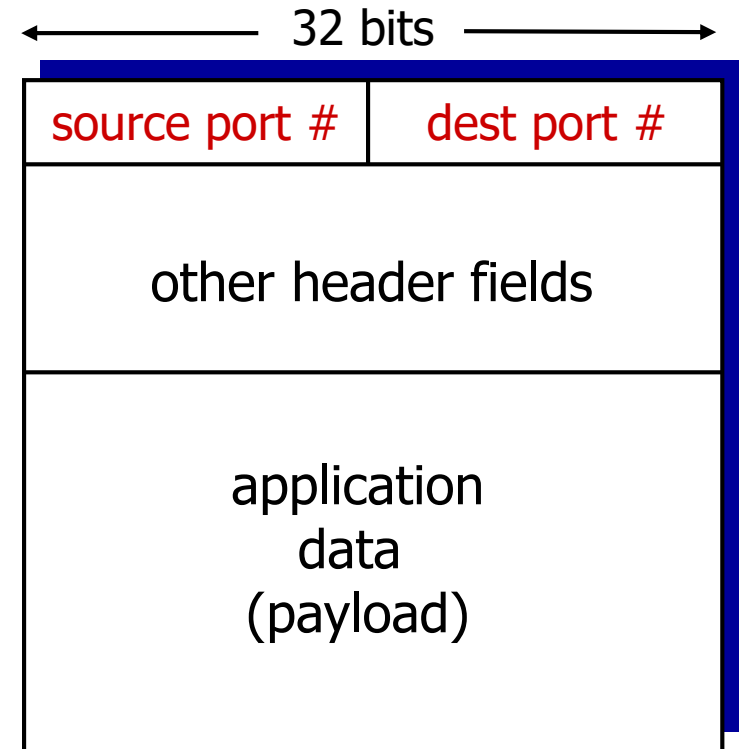
demultiplexing at receiver:

use header info to deliver received segments to correct socket



How demultiplexing works

- host receives IP datagrams
 - each datagram has source IP address, destination IP address
 - each datagram carries one transport-layer segment
 - each segment has source, destination port number
- host uses *IP addresses & port numbers* to direct segment to appropriate socket



TCP/UDP segment format

Connectionless demultiplexing

- *recall*: created socket has host-local port #:

```
DatagramSocket mySocket1  
= new DatagramSocket(12534) ;
```

- *recall*: when creating datagram to send into UDP socket, must specify

- destination IP address
- destination port #

-
- when host receives UDP segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #



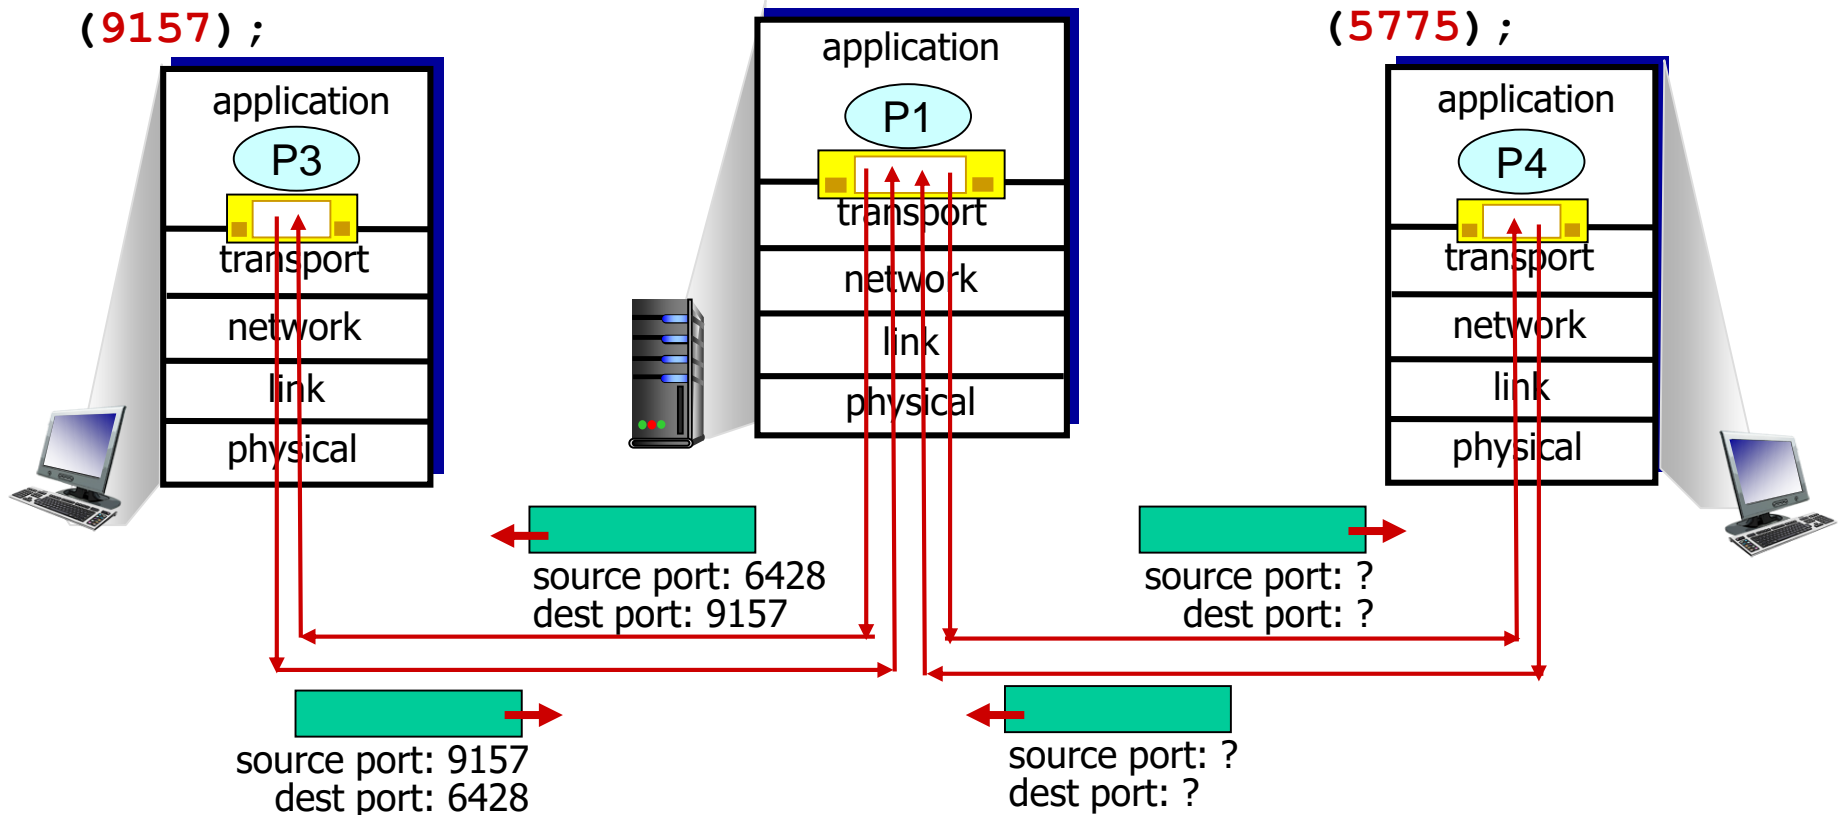
IP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at dest

Connectionless demux: example

```
DatagramSocket  
mySocket2 = new  
DatagramSocket  
(9157);
```

```
DatagramSocket  
serverSocket = new  
DatagramSocket  
(6428);
```

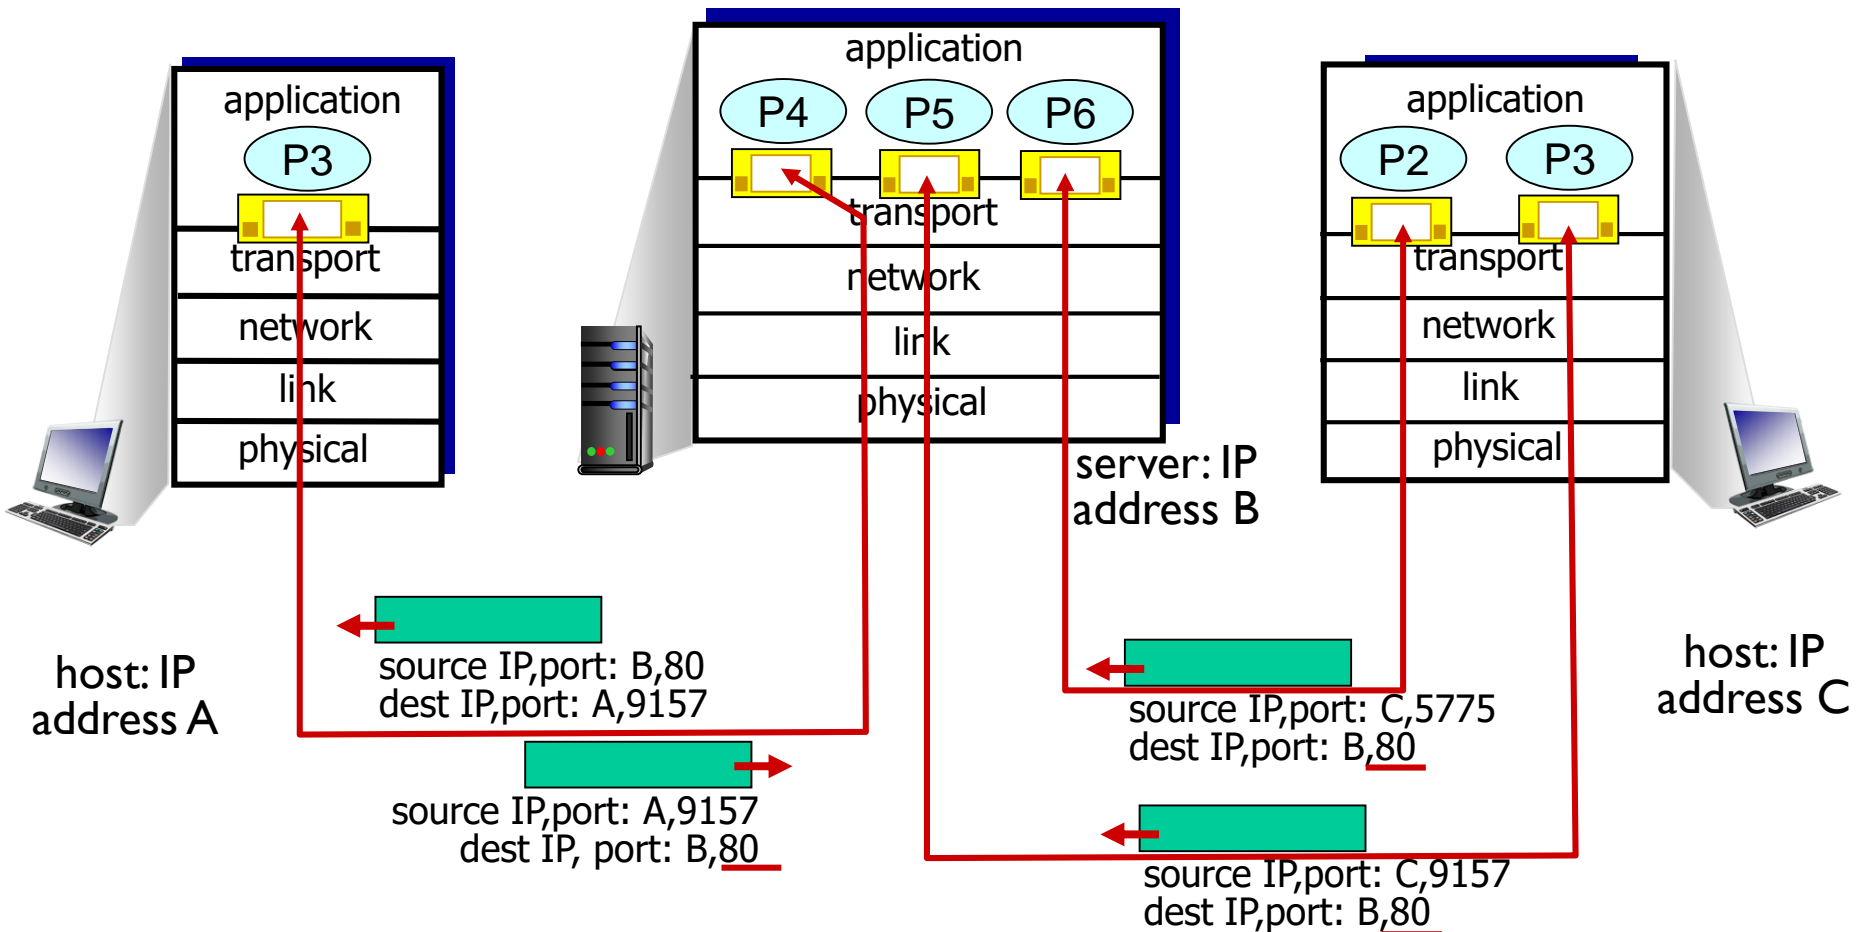
```
DatagramSocket  
mySocket1 = new  
DatagramSocket  
(5775);
```



Connection-oriented demux

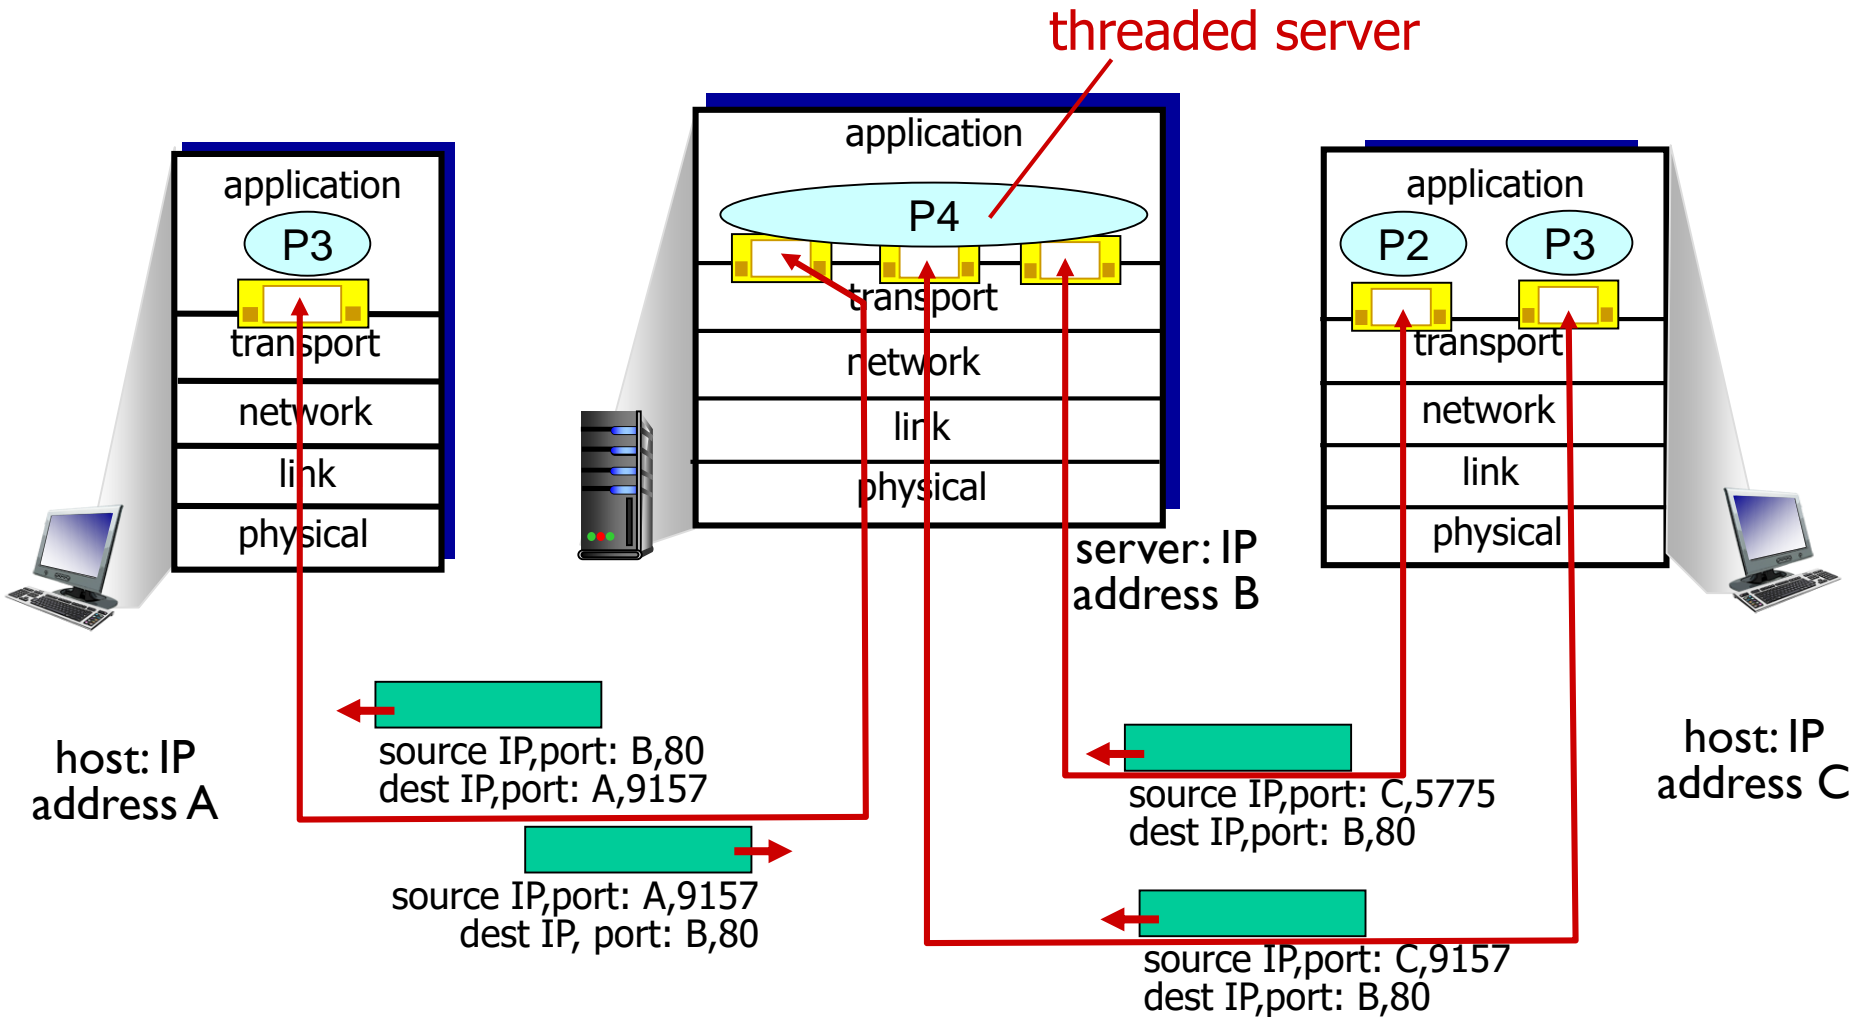
- TCP socket identified by 4-tuple:
 - source IP address
 - source port number
 - dest IP address
 - dest port number
- demux: receiver uses all four values to direct segment to appropriate socket
- server host may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
- web servers have different sockets for each connecting client
 - non-persistent HTTP will have different socket for each request

Connection-oriented demux: example



three segments, all destined to IP address: B,
dest port: 80 are demultiplexed to *different* sockets

Connection-oriented demux: example



Typical port numbers for applications

- Port 80: Standard HTTP (Hypertext Transfer Protocol) – browse the web
- Port 22: ssh (secure shell) log in from remote computer
- Port 25: SMTP (simple mail transfer protocol) – send mail via this host
- Port 143: IMAP (Internet Message Access Protocol) – read your email
- Port 443: HTTPS (secure HTTP) browse the web securely
- These are just conventions, you could change them if you so wished on your own host.

Transport layer outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

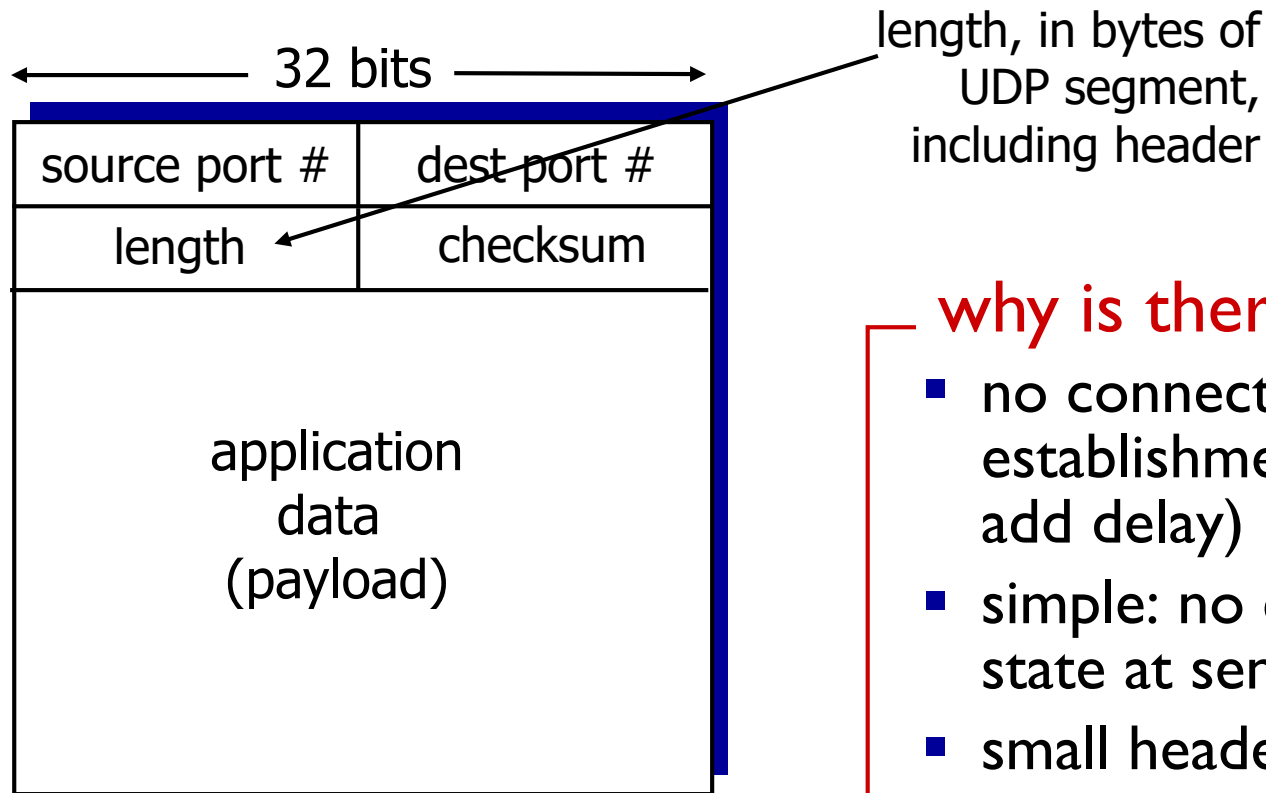
3.6 principles of congestion control

3.7 TCP congestion control

UDP: User Datagram Protocol [RFC 768]

- Simplest usable Internet transport protocol
- “best effort” service, UDP segments may be:
 - lost
 - delivered out-of-order to app
- *connectionless*:
 - no “handshaking” between UDP sender, receiver (can send immediately without asking first)
 - each UDP packet handled independently of others
- UDP use:
 - streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS
 - SNMP
- reliable transfer over UDP:
 - add reliability at application layer
 - application-specific error recovery!

UDP: segment header



UDP segment format

why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control: UDP can blast away as fast as desired

UDP checksum

Goal: detect “errors” (e.g., flipped bits – 0 becomes 1 or 1 becomes 0) in transmitted segment


sender:

- treat segment contents, including header fields, as sequence of 16-bit integers
- checksum: addition (one's complement sum) of segment contents
- sender puts checksum value into UDP checksum field

receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected.
But maybe errors nonetheless? More later
....

UDP Encapsulation/decapsulation




What have we learned?

- Brief look at Application Layer (layer seven) protocols which programmers use.
- Transport Layer introduction (layer four).
 - Deliver data to application sockets themselves.
 - Responsible for “end-to-end” connection.
- User Datagram Protocol (transport layer):
 - Simple, connectionless.
 - Packets may be lost and reordered.
 - Useful for simple situations or real-time.

Structure of course

- Week 1
 - Introduction to IP Networks
 - The Transport layer (part I)
- Week 2
 - The Transport layer (part II)
 - The Network layer (part I)
 - Class test
- Week 3
 - The Network layer (part II)
 - The Data link layer (part I)
 - Router lab tutorial (assessed lab work after this week)
- Week 4
 - The Data link layer (part II)
 - Network management and security
 - Class test

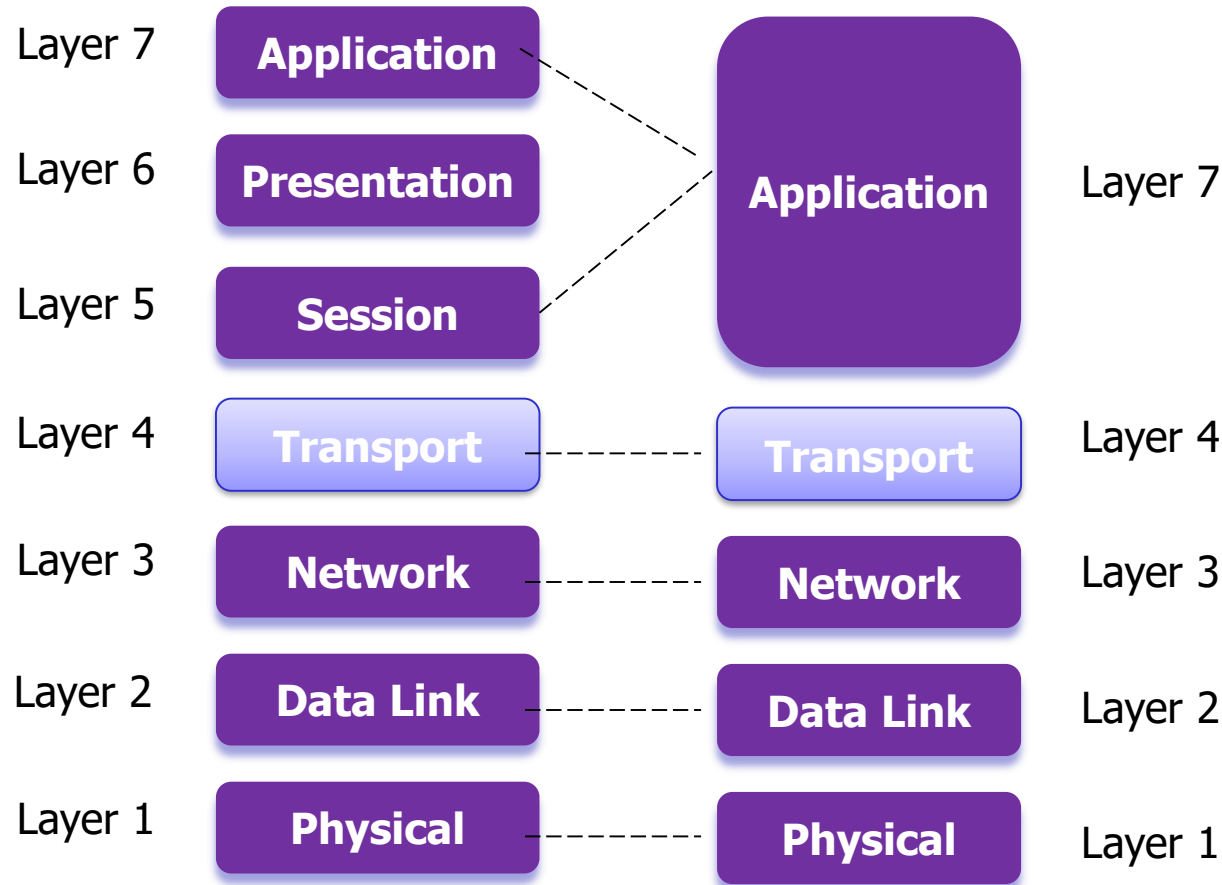
ISO/OSI (left) vs TCP/IP (right)



How to remember the layers

- Please Do Not Throw Sausage + Pizza Away
- **P**lease – **P**hysical
- **D**o – **D**atalink
- **N**ot – **N**etwork
- **T**hrow – **T**ransport
- **S**ausage – **S**ession
- **P**izza – **P**resentation
- **A**way – **A**pplication

Transport Layer



Week 1: Transport Layer (part II)

Our goals:

- understand principles behind transport layer services:
 - multiplexing, demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
- learn about Internet transport layer protocols:
 - UDP: connectionless transport
 - TCP: connection-oriented reliable transport
 - TCP congestion control

Transport layer outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP


- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Principles of reliable data transfer

- important in application, transport, link layers
 - top-10 list of important networking topics!




(a) provided service

- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Principles of reliable data transfer


- important in application, transport, link layers
 - top-10 list of important networking topics!



- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Principles of reliable data transfer

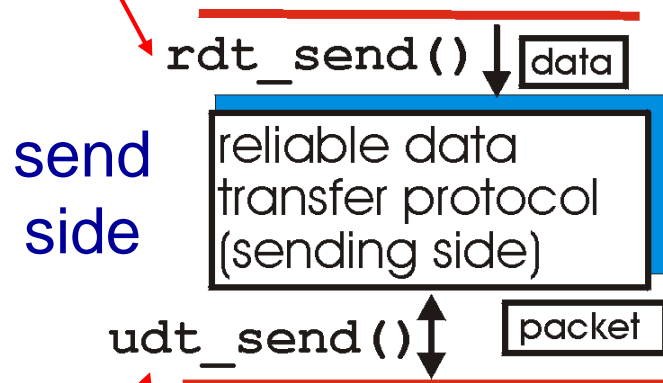
- important in application, transport, link layers
 - top-10 list of important networking topics!



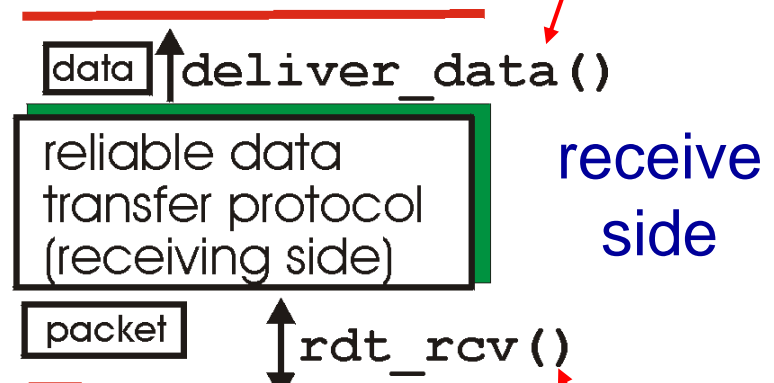
- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Reliable data transfer: getting started

rdt_send() : called from above,
(e.g., by app.). Passed data to
deliver to receiver upper layer



deliver_data() : called by
rdt to deliver data to upper




udt_send() : called by rdt,
to transfer packet over
unreliable channel to receiver

rdt_rcv() : called when packet
arrives on rcv-side of channel

What is a Finite State Machine

- A finite state machine shows how a system works with states and transitions between them.
- Consider a coin operated turnstile – we put in a coin, it allows us in.
- It has two states "locked" "unlocked".
- It has two events "push" and "insert coin"



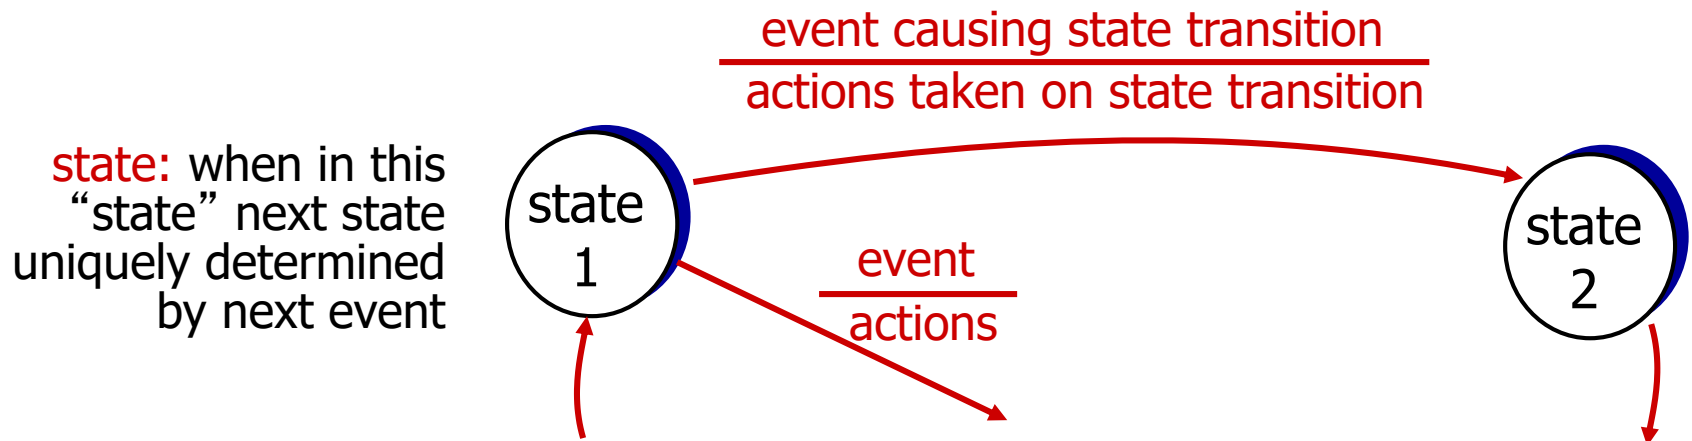
Diagrams from wikipedia



Reliable data transfer: getting started


We'll:

- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer
 - but control info will flow on both directions!
- use finite state machines (FSM) to specify sender, receiver




rdt1.0: reliable transfer over a reliable channel

- underlying channel perfectly reliable
 - no bit errors
 - no loss of packets
- separate FSMs for sender, receiver:
 - sender sends data into underlying channel
 - receiver reads data from underlying channel



sender




receiver

rdt2.0: channel with bit errors

- We drop the assumption that the channel has no errors.
- Now the channel could have bit errors.
- We need to be able to detect and recover from errors in transmission.
- Consider this question:
 - How do “humans” recover from errors in conversation?

Reliable data transfer: getting started

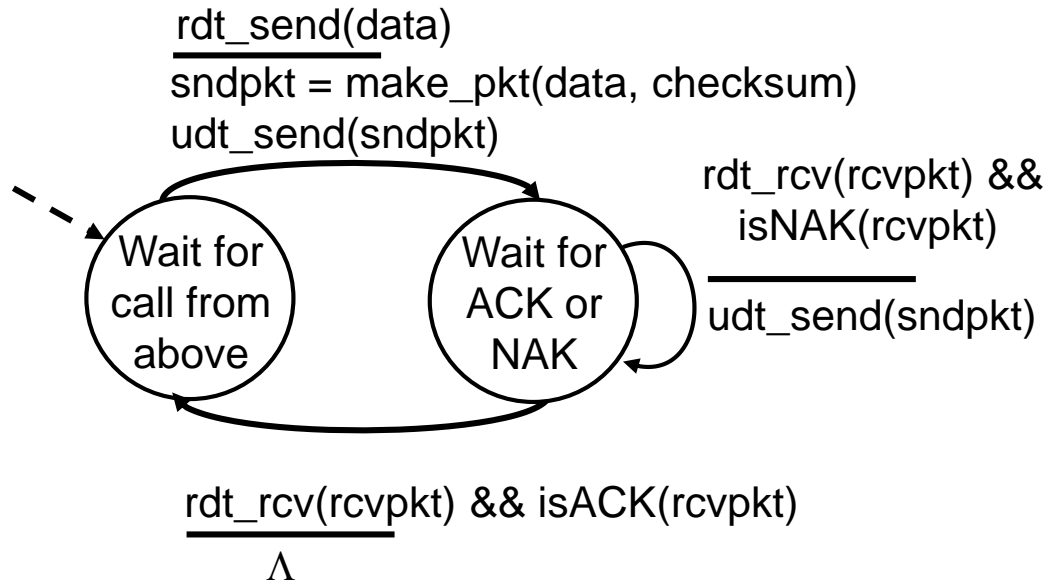
- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer
 - but control info will flow on both directions!
- use finite state machines (FSM) to specify sender, receiver



rdt2.0: channel with bit errors

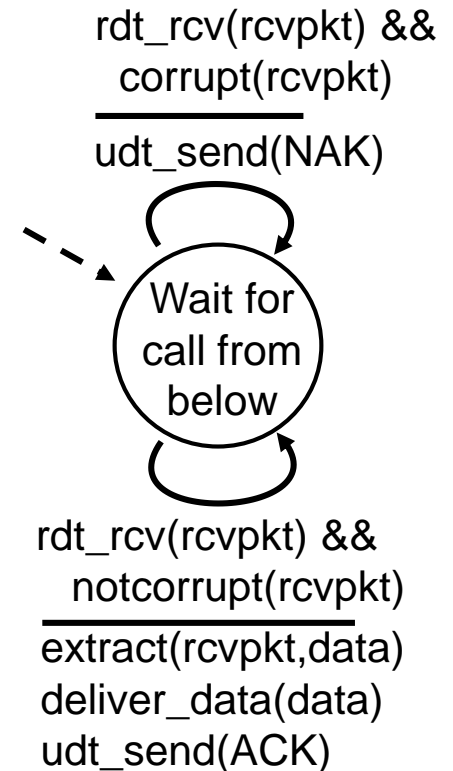
- underlying channel may flip bits in packet
 - checksum to detect bit errors
- *the question*: how to recover from errors:
 - *acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
 - *negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
 - sender retransmits pkt on receipt of NAK
- new mechanisms in `rdt2.0` (beyond `rdt1.0`):
 - error detection
 - feedback: control msgs (ACK,NAK) from receiver to sender

rdt2.0: FSM specification

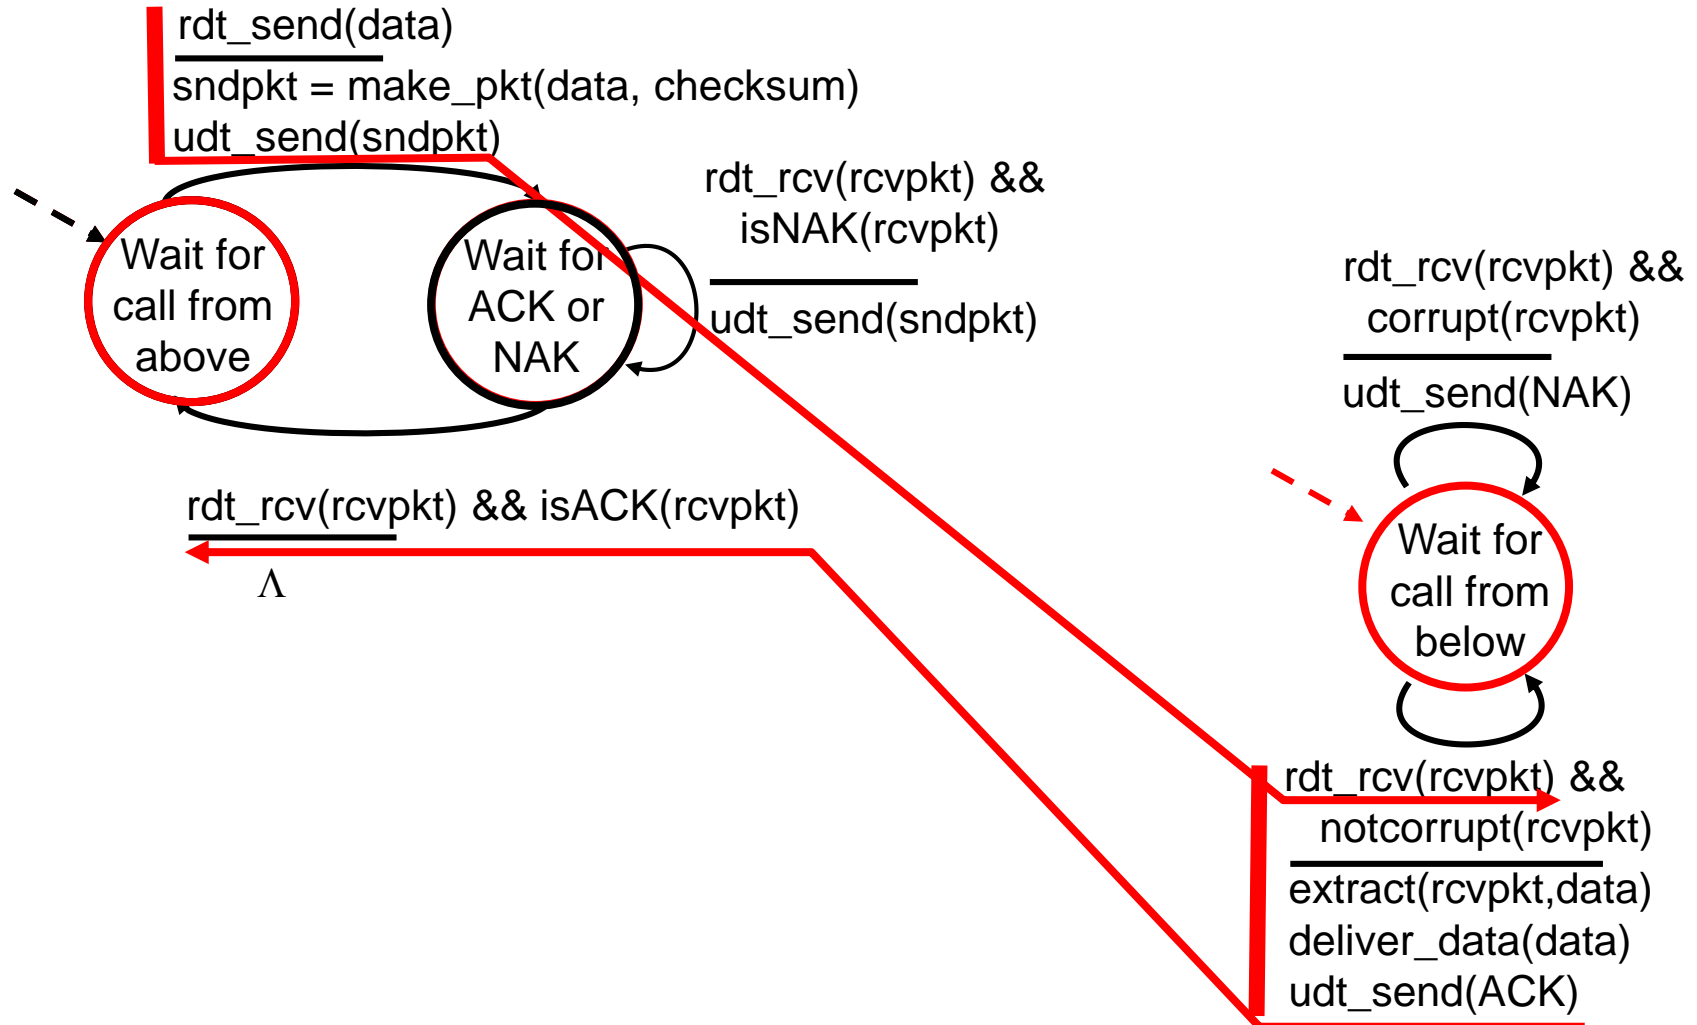


sender

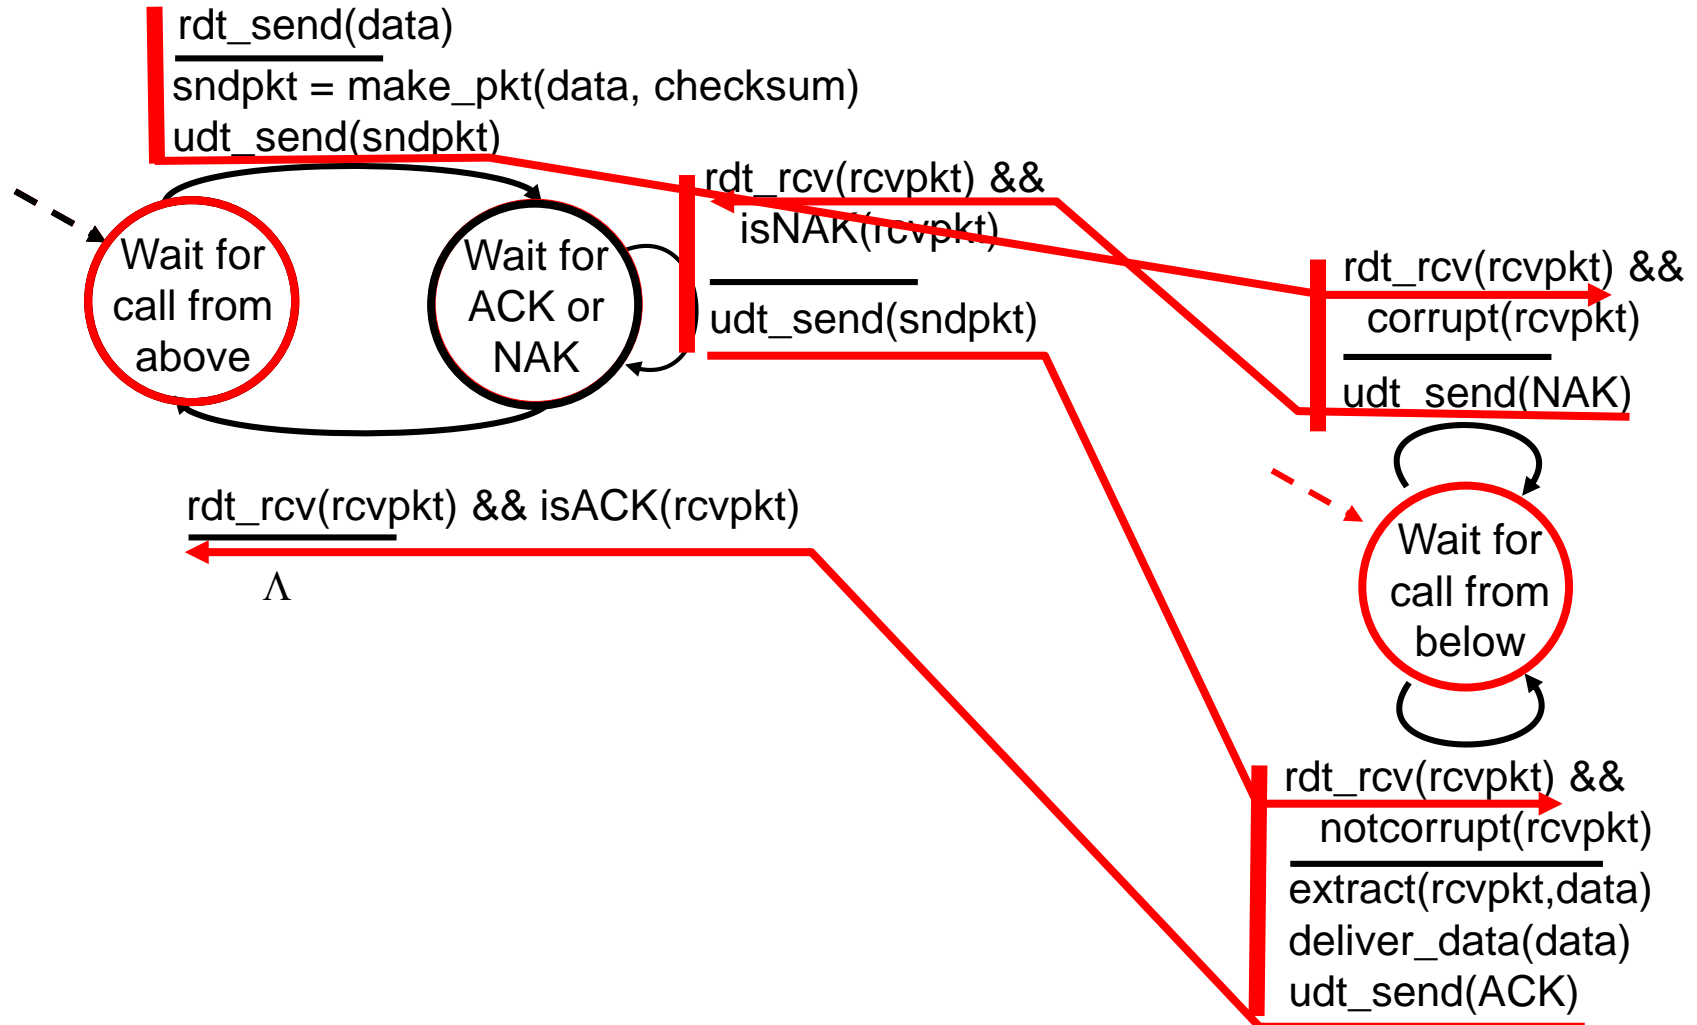
receiver



rdt2.0: operation with no errors



rdt2.0: error scenario



rdt2.0 has a fatal flaw!

what happens if ACK/NAK corrupted?

- sender doesn't know what happened at receiver!
- can't just retransmit:
possible duplicate

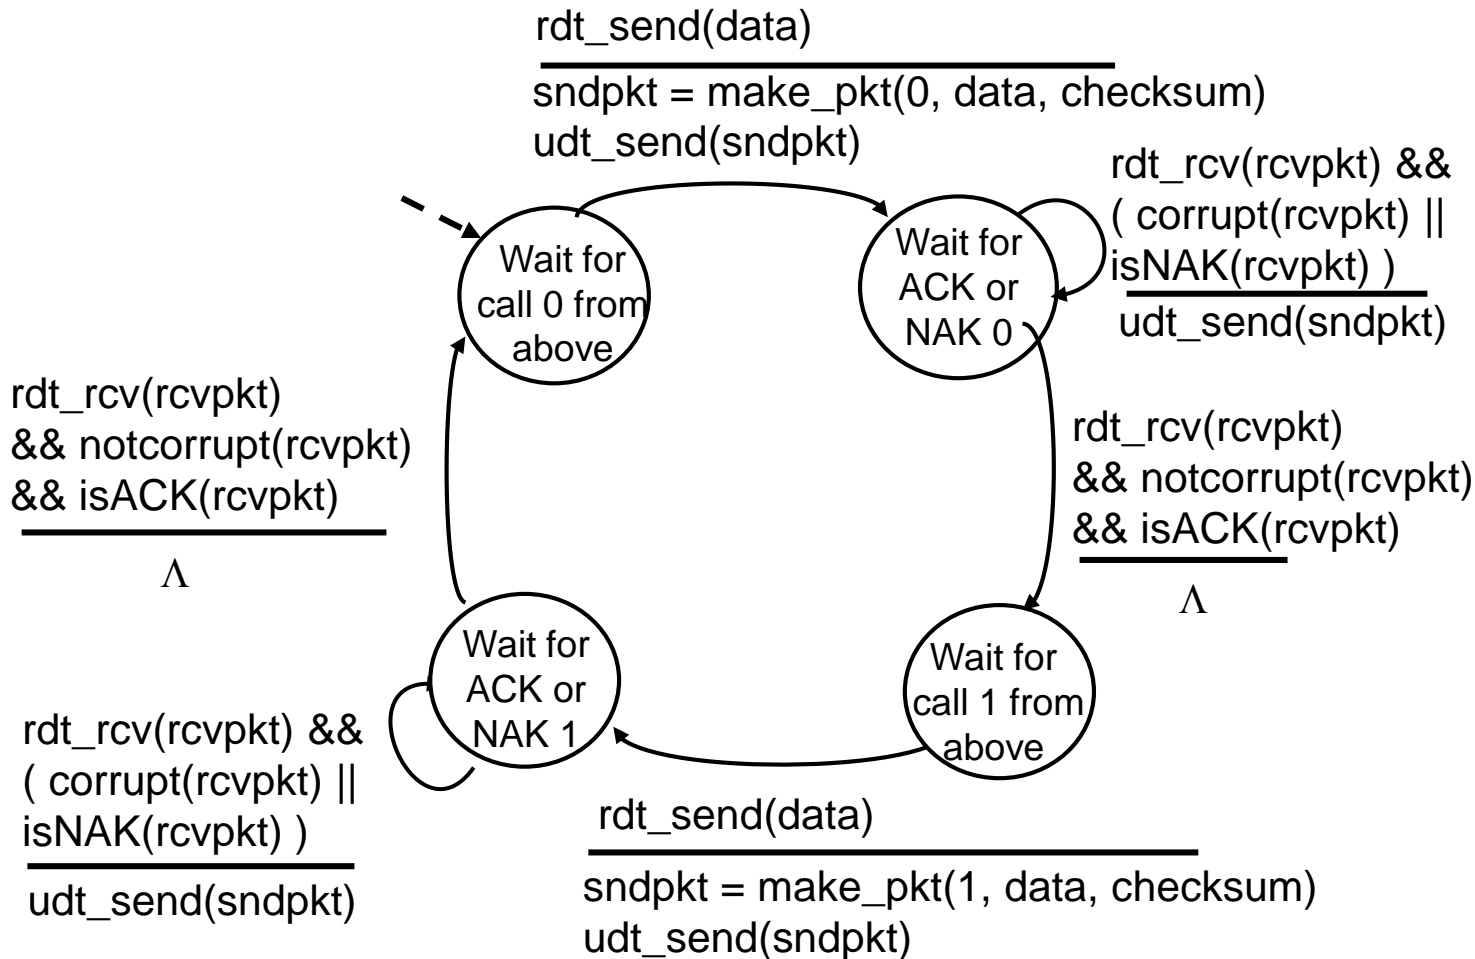
handling duplicates:

- sender retransmits current pkt if ACK/NAK corrupted
- sender adds *sequence number* (seq) to each pkt
- receiver discards (doesn't deliver up) duplicate pkt


stop and wait

sender sends one packet,
then waits for receiver
response

rdt2.1: sender, handles corrupt ACK/NAKs



rdt2.1: receiver, handles corrupt ACK/NAKs



rdt2.1: discussion

sender:

- seq # added to pkt
- two seq. #'s (0,1) will suffice. Why?
- must check if received ACK/NAK corrupted
- twice as many states
 - state must “remember” whether “expected” pkt should have seq # of 0 or 1

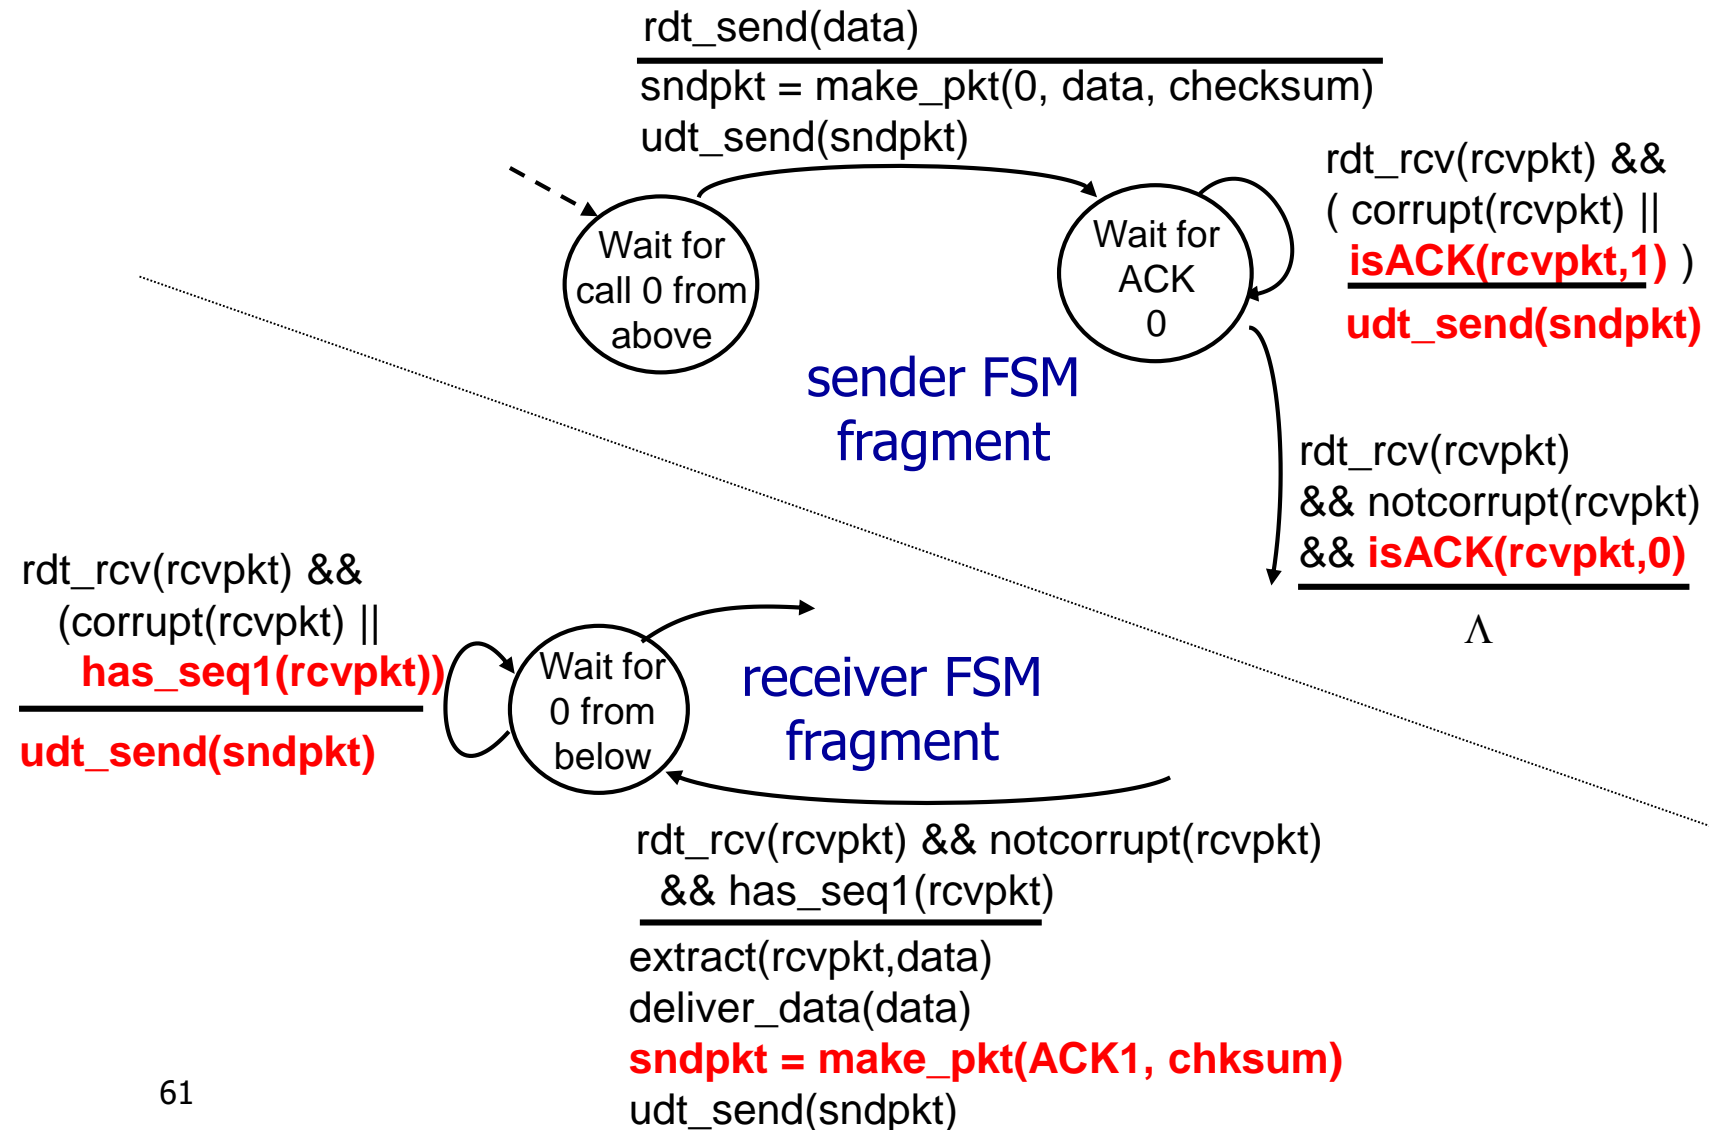
receiver:

- must check if received packet is duplicate
 - state indicates whether 0 or 1 is expected pkt seq #
- note: receiver can *not* know if its last ACK/NAK received OK at sender

rdt2.2: a NAK-free protocol

- same functionality as rdt2.1, using ACKs only
- instead of NAK, receiver sends ACK for last pkt received OK
 - receiver must *explicitly* include seq # of pkt being ACKed
- duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

rdt2.2: sender, receiver fragments



rdt3.0: channels with errors *and* loss

new assumption:


underlying channel can also lose packets (data, ACKs)

- checksum, seq. #, ACKs, retransmissions will be of help ... but not enough

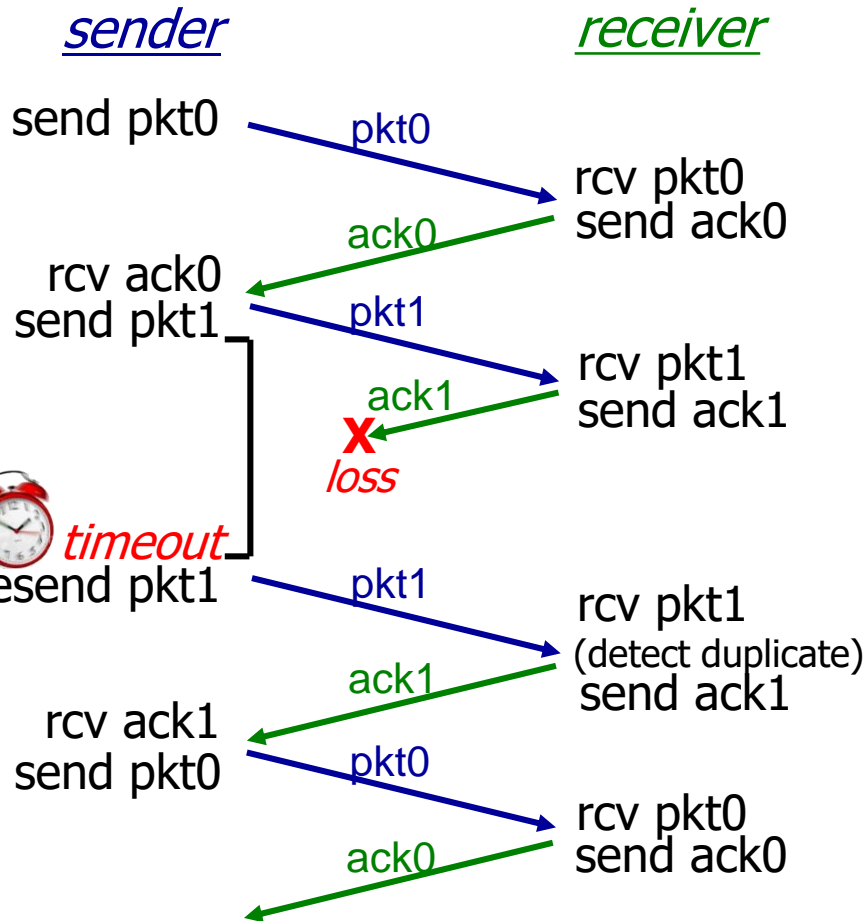
approach: sender waits “reasonable” amount of time for ACK

- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but seq. #'s already handles this
 - receiver must specify seq # of pkt being ACKed
- requires countdown timer


rdt3.0 sender



rdt3.0 in action



(c) ACK loss



(d) premature timeout/ delayed ACK

What have we learned?

- We are starting on our journey to create “reliable data transfer” over an unreliable network.
- Key concepts:
 - Sequence number – measures which packet we are currently sending
 - ACK (acknowledgement) – numbered, says that a specific packet has definitely been arrived
 - Timeout – a time to wait before we think a packet is lost

Transport layer summary

- So far:
 - 3.1 transport-layer services
 - 3.2 multiplexing and demultiplexing
 - 3.3 connectionless transport: UDP
 - 3.4 principles of reliable data transfer
- Next week of lectures:
 - 3.4 principles of reliable data transfer (continued)
 - 3.5 connection-oriented transport: TCP
 - 3.6 principles of congestion control
 - 3.7 TCP congestion control